

NEURAL NETWORKS

JAMIE CLINE

- What is a neural network?
- How do they learn?
- The mathematics involved

Fig: gradient descent

CONTENTS

- What is a neural network?
- How do they learn?
- Backpropagation
- Weight Decay
- The Hessian
- Modern Optimizers
- Conclusion



Fig: gradient descent


Intro to Neurons & Layers

- Neural networks are functions made of neurons and layers.
- Each neuron stores its activation
- A non-linear activation function is applied

$$z^{(\ell+1)} = \mathbf{W}^{(\ell+1)} a^{(\ell)} + b^{(\ell+1)}$$

$$a^{(\ell+1)} = \sigma(z^{(\ell+1)})$$

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$




How do they learn?

- The weights, biases, and learning rate tunes the model.
- The cost function tells the network how good the parameters are.
- Gradient descent is used to minimize C

$$C_n = \frac{1}{n} \sum_j (a_j^{(\ell)} - y_j)^2$$

params := *params* – $\eta \nabla C$



Backpropagation


- How do we know how much each weight effects the final error?
- Output layer requests change in previous layer
- Requests propagate backwards

$$\frac{\partial C}{\partial a^{(\ell)}} = 2(a^{(\ell)} - y)$$

$$\frac{\partial a^{(\ell)}}{\partial z^{(\ell)}} = \sigma'(z^{(\ell)})$$

$$\frac{\partial z^{(\ell)}}{\partial w^{(\ell)}} = a^{(\ell-1)}$$

$$\frac{\partial C}{\partial w^{(\ell)}} = 2(a^{(\ell)} - y)\sigma'(z^{(\ell)})a^{(\ell-1)}$$



Weight Decay


- Adjusting parameters like this can cause overfitting
- Adding regularization to the cost can reduce this.

$$C_{total} = C_{data} + \lambda R(W)$$

$$\left\| W^l \right\|_F^2 = \sum_{i,j} \left(W_{ij}^{(l)} \right)^2$$

$$R(W) = \sum_l \left\| W^{(l)} \right\|_F^2$$

$$W^{(l)} := W^{(l)} - \eta \left(\frac{\partial C_{data}}{\partial W^{(l)}} + 2\lambda W^{(l)} \right)$$




The Hessian



- The hessian is a matrix of second derivatives.
- We can approximate cost with Taylor expansion up to second order close to θ .

$$g(\theta) = \nabla_{\theta} C(\theta)$$

$$H(\theta) = \nabla_{\theta}^2 C(\theta)$$

$$C(\theta + \Delta) \approx C(\theta) + g(\theta)^T \Delta + \frac{1}{2} \Delta^T H(\theta) \Delta$$





Adam (AME) Optimisation

- Incorporates momentum into gradient descent
- Keeps track of two running averages
- Direction moment :


$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$$

- Scale of curvature moment:

$$\nu_t = \beta_2 \nu_{t-1} + (1 - \beta_2) g_t^2$$


- And then updates each parameter:

$$\theta_{t+1} = \theta_t - \alpha \frac{m_t}{\sqrt{\nu_t} + \epsilon}$$



CONCLUSION

- Neural networks can approximate any function
- Layers and activations give universal expressivity
- Training adjusts weights to minimize cost
- Regularization and optimization shape performance
- Foundation of modern AI and scientific computing



```
main.py > ...
1  import numpy as np
2  import matplotlib.pyplot as plt
3  from mpl_toolkits.mplot3d import Axes3D
4
5  # Create a grid
6  x = np.linspace(-3, 3, 400)
7  y = np.linspace(-3, 3, 400)
8  X, Y = np.meshgrid(x, y)
9  R = X**2 + Y**2
10
11 # Cost function
12 Z = np.sin(R)+(R) + 2*np.sin(3*X)*np.sin(3*Y)
13
14 # Plot
15 fig = plt.figure(figsize=(8, 6))
16 ax = fig.add_subplot(111, projection='3d')
17 ax.plot_surface(X, Y, Z, cmap='plasma', rstride=3, cstride=3, linewidth=0, alpha=0.9)
18
19 # Styling
20 ax.set_xticks([])
21 ax.set_yticks([])
22 ax.set_zticks([])
23 ax.set_box_aspect((1,1,0.5))
24 plt.tight_layout()
25 plt.show()
```